


```

75%    3.677083    12.500000    18.100000    0.000000    0.624000    6.623500
      94.075000    5.188425    24.000000    666.000000    20.200000    396.225000
      16.955000    25.000000
max    88.976200    100.000000    27.740000    1.000000    0.871000    8.780000
      100.000000    12.126500    24.000000    711.000000    22.000000    396.900000
      37.970000    50.000000

```

#View dataset features and attribute information

df.head()

```

      rim    zn    indus  chas  nox    rm    age    dis    rad    tax    ptratio  black
lstat  medv
0    0.00632    18.0    2.31    0    0.538  6.575  65.2    4.0900  1    296.0    15.3
      396.90  4.98    24.0
1    0.02731    0.0    7.07    0    0.469  6.421  78.9    4.9671  2    242.0    17.8
      396.90  9.14    21.6
2    0.02729    0.0    7.07    0    0.469  7.185  61.1    4.9671  2    242.0    17.8
      392.83  4.03    34.7
3    0.03237    0.0    2.18    0    0.458  6.998  45.8    6.0622  3    222.0    18.7
      394.63  2.94    33.4
4    0.06905    0.0    2.18    0    0.458  7.147  54.2    6.0622  3    222.0    18.7
      396.90  5.33    36.2

```

Assign feature variable to df_x

df_x=df

print(df_x)

```

crim  zn  indus  chas  nox    rm  age    dis  rad  tax \
0  0.00632  18.0  2.31    0  0.538  6.575  65.2  4.0900  1  296.0
1  0.02731  0.0  7.07    0  0.469  6.421  78.9  4.9671  2  242.0
2  0.02729  0.0  7.07    0  0.469  7.185  61.1  4.9671  2  242.0
3  0.03237  0.0  2.18    0  0.458  6.998  45.8  6.0622  3  222.0
4  0.06905  0.0  2.18    0  0.458  7.147  54.2  6.0622  3  222.0
..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..
501  0.06263  0.0  11.93    0  0.573  6.593  69.1  2.4786  1  273.0
502  0.04527  0.0  11.93    0  0.573  6.120  76.7  2.2875  1  273.0
503  0.06076  0.0  11.93    0  0.573  6.976  91.0  2.1675  1  273.0
504  0.10959  0.0  11.93    0  0.573  6.794  89.3  2.3889  1  273.0
505  0.04741  0.0  11.93    0  0.573  6.030  80.8  2.5050  1  273.0

```

```

      ptratio  black  lstat  medv
0    15.3  396.90  4.98  24.0
1    17.8  396.90  9.14  21.6
2    17.8  392.83  4.03  34.7
3    18.7  394.63  2.94  33.4
4    18.7  396.90  5.33  36.2

```

```

..      ...      ...      ...      ...
501    21.0 391.99 9.67 22.4
502    21.0 396.90 9.08 20.6
503    21.0 396.90 5.64 23.9
504    21.0 393.45 6.48 22.0
505    21.0 396.90 7.88 11.9

```

```
# Assign medv variable to df_y
```

```
df_y=df.medv
```

```
print(df_y)
```

```
0    24.0
```

```
1    21.6
```

```
2    34.7
```

```
3    33.4
```

```
4    36.2
```

```
...
```

```
501    22.4
```

```
502    20.6
```

```
503    23.9
```

```
504    22.0
```

```
505    11.9
```

```
Name: medv, Length: 506, dtype: float64
```

```
#Initialize the linear regression model
```

```
reg =linear_model.LinearRegression()
```

```
#Split the data into 67% training and 33% testing data
```

```
#NOTE: We have to split the dependent variables (x) and the target or independent variable (y)
```

```
x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.33, random_state=42)
```

```
#Train our model with the training data
```

```
reg.fit(x_train, y_train)
```

```
#Print the coefecients/weights for each feature/column of our model
```

```
print(reg.coef_)
```

```
[ 2.03880984e-16  3.74700271e-16 -8.77986919e-16  1.44095737e-14
 -1.37739992e-14 -1.45813673e-15 -3.47378376e-16 -6.97846728e-16
  6.24500451e-16 -3.81639165e-17  3.95462742e-16  2.94902991e-17
  1.08105799e-15  1.00000000e+00]
```

```
#print our price predictions on our test data
```

```
y_pred = reg.predict(x_test)
```

```
print(y_pred)
```

```
23.6 32.4 13.6 22.8 16.1 20. 17.8 14. 19.6 16.8 21.5 18.9 7. 21.2
```

18.5 29.8 18.8 10.2 50. 14.1 25.2 29.1 12.7 22.4 14.2 13.8 20.3 14.9
21.7 18.3 23.1 23.8 15. 20.8 19.1 19.4 34.7 19.5 24.4 23.4 19.7 28.2
50. 17.4 22.6 15.1 13.1 24.2 19.9 24. 18.9 35.4 15.2 26.5 43.5 21.2
18.4 28.5 23.9 18.5 25. 35.4 31.5 20.2 24.1 20. 13.1 24.8 30.8 12.7
20. 23.7 10.8 20.6 20.8 5. 20.1 48.5 10.9 7. 20.9 17.2 20.9 9.7
19.4 29. 16.4 25. 25. 17.1 23.2 10.4 19.6 17.2 27.5 23. 50. 17.9
9.6 17.2 22.5 21.4 12. 19.9 19.4 13.4 18.2 24.6 21.1 24.7 8.7 27.5
20.7 36.2 31.6 11.7 39.8 13.9 21.8 23.7 17.6 24.4 8.8 19.2 25.3 20.4
23.1 37.9 15.6 45.4 15.7 22.6 14.5 18.7 17.8 16.1 20.6 31.6 29.1 15.6
17.5 22.5 19.4 19.3 8.5 20.6 17. 17.1 14.5 50. 14.3 12.6 28.7 21.2
19.3 23.1 19.1 25. 33.4 5. 29.6 18.7 21.7 23.1 22.8 21. 48.8]

```
#print the predicted price and actual price of houses from the testing data set row 0
y_pred[0]
23.599999999999994
```

```
y_test[0]
24.0
```

Conclusion:

By looking at the predicted values that the model came up with, and the actual values of the testing data set, it looks like the model is pretty good at making predictions. It's not exact, but it is pretty close and much better than guessing.

To check the model's performance/accuracy I will use a metric called mean squared error (MSE)

```
print(np.mean((y_pred-y_test)**2))
2.967534119029605e-28
```

Conclusion:

The mean squared error function returns a value of about 20.72. This is really good! If our model had predicted the exact value as the actual values then the mean squared error function would have returned a value of 0.